

# Algorithm Design & Analysis (CS3383)<sup>1</sup>

---

## Unit 5 : Dynamic Multithreaded Algorithms

Rasoul Shahsavarifar

April 1, 2019

---

<sup>1</sup>Thanks to Dr. Patricia Evans and Dr. David Bremner at UNB for sharing the teaching stuffs 

# Outline

## Dynamic Multithreaded Algorithms Fork-Join Model

## Using More than One Processor

- Capabilities
- Span, Work, And Parallelism
- Parallel Loops
- Scheduling
- Race Conditions

# Contents

## Dynamic Multithreaded Algorithms Fork-Join Model

### Using More than One Processor

Capabilities

Span, Work, And Parallelism

Parallel Loops

Scheduling

Race Conditions

# Introduction to Parallel Algorithms

## Dynamic Multithreading

- ▶ Also known as the *fork-join* model
- ▶ Shared memory, *multicore*
- ▶ Cormen et. al 3rd edition, Chapter 27

# Introduction to Parallel Algorithms

## Dynamic Multithreading

- ▶ Also known as the *fork-join* model
- ▶ Shared memory, *multicore*
- ▶ Cormen et. al 3rd edition, Chapter 27

# Multithreading

So far we've been looking at traditional *serial algorithmics*, designing algorithms to run on a single processor and analyzing the single-processor running time.

However, it's quite common for computers to have multiprocessors, and for code to be multithreaded.

At a larger parallel scale, processors may have access to many processors.

So, how does this change what we do for algorithm design and analysis?

# Contents

Dynamic Multithreaded Algorithms  
Fork-Join Model

Using More than One Processor  
Capabilities  
Span, Work, And Parallelism  
Parallel Loops  
Scheduling  
Race Conditions

# Concurrency Capabilities

We examine a shared-memory framework, so multiple processors can access the same memory. This is consistent with many architectures, and also eliminates passing data between processors.

The model we use is often referred to as the *fork-join* model. The structures we seek to parallelize are the two fundamental code structures: the *branch* and the *loop*. While loops don't parallelize well due to their structure, so the loops are limited to iterative *for* loops with a counter.

New keywords:

**parallel** run the loop potentially concurrently

**spawn** run the called routine potentially concurrently

**sync** wait for all spawned children to complete

These keywords can be added to serial code, and can also be removed from parallel code to serialize it.



# Writing parallel (pseudo)-code

## Keywords

`parallel` Run the loop (potentially) concurrently

`spawn` Run the procedure (potentially) concurrently

`sync` Wait for all spawned children to complete.

# Writing parallel (pseudo)-code

## Keywords

- `parallel` Run the loop (potentially) concurrently
- `spawn` Run the procedure (potentially) concurrently
- `sync` Wait for all spawned children to complete.

## Serialization

- ▶ remove keywords from parallel code yields correct serial code
- ▶ Adding parallel keywords to correct serial code might break it

# Writing parallel (pseudo)-code

## Keywords

- `parallel` Run the loop (potentially) concurrently
- `spawn` Run the procedure (potentially) concurrently
- `sync` Wait for all spawned children to complete.

## Serialization

- ▶ remove keywords from parallel code yields correct serial code
- ▶ Adding parallel keywords to correct serial code might break it
  - ▶ missing sync

# Writing parallel (pseudo)-code

## Keywords

- `parallel` Run the loop (potentially) concurrently
- `spawn` Run the procedure (potentially) concurrently
- `sync` Wait for all spawned children to complete.

## Serialization

- ▶ remove keywords from parallel code yields correct serial code
- ▶ Adding parallel keywords to correct serial code might break it
  - ▶ missing sync
  - ▶ loop iterations not independent

# Fibonacci Example

```
function FIB( $n$ )  
  if  $n \leq 1$  then  
    return  $n$   
  else  
     $x = \text{Fib}(n - 1)$   
     $y = \text{Fib}(n - 2)$   
  
    return  $x + y$   
  end if  
end function
```

# Fibonacci Example

```
function FIB( $n$ )  
  if  $n \leq 1$  then  
    return  $n$   
  else  
     $x =$  spawn Fib( $n - 1$ )  
     $y =$  Fib( $n - 2$ )  
    sync  
    return  $x + y$   
  end if  
end function
```

- Code in C, Java, Clojure and Racket available from <http://www.cs.unb.ca/~bremner/teaching/cs3383/examples>

# Contents

## Dynamic Multithreaded Algorithms

Fork-Join Model

## Using More than One Processor

Capabilities

Span, Work, And Parallelism

Parallel Loops

Scheduling

Race Conditions

# Computation DAG

# Strands

Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.

**function** FIB( $n$ )

**if  $n \leq 1$  then**

```
return  $n$ 
```

**else**


$$x = \text{spawn Fib}(n - 1)$$
$$y = \text{Fib}(n - 2)$$

sync

```
return  $x + y$ 
```

**end if**

end function






# Computation DAG

## Strands

Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.

nodes strands  
down edges spawn




# Computation DAG

## Strands

Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.

nodes strands  
down edges spawn  
up edges return



# Computation DAG

## Strands


Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.

nodes strands

down edges spawn

up edges return

horizontal edges sequential



# Computation DAG

## Strands

Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.


nodes strands

down edges spawn

up edges return

horizontal edges sequential

critical path longest path in DAG



# Computation DAG

## Strands

Seq. inst. with no *parallel*, *spawn*, return from *spawn*, or *sync*.

nodes strands


down edges spawn

up edges return

horizontal edges sequential

critical path longest path in DAG

span weighted length of  
critical path  $\equiv$  lower  
bound on time



# Work and Speedup

$T_1$  *Work*, sequential time.

# Work and Speedup

$T_1$  *Work*, sequential time.

$T_p$  Time on  $p$  processors.

# Work and Speedup

$T_1$  Work, sequential time.

$T_p$  Time on  $p$  processors.

## Work Law

$$T_p \geq T_1/p$$

$$\text{speedup} := T_1/T_p \leq p$$



# Parallelism

$T_p$  Time on  $p$  processors.

# Parallelism

We could idle processors:

$$T_p \geq T_\infty \quad (1)$$

$T_p$  Time on  $p$  processors.

$T_\infty$  *Span*, time given  
unlimited processors.

# Parallelism

We could idle processors:

$$T_p \geq T_\infty \quad (1)$$

$T_p$  Time on  $p$  processors.  
 $T_\infty$  *Span*, time given  
unlimited processors.


Best possible speedup:

$$\begin{aligned} \text{parallelism} &= T_1/T_\infty \\ &\geq T_1/T_p = \text{speedup} \end{aligned}$$

# Span and Parallelism Example

Assume strands are unit cost.

►  $T_1 = 17$




# Span and Parallelism Example

Assume strands are unit cost.

►  $T_1 = 17$


►  $T_\infty = 8$




# Span and Parallelism Example

Assume strands are unit cost.

- ▶  $T_1 = 17$
- ▶  $T_\infty = 8$
- ▶ Parallelism = 2.125 for **this** input size.




## Composing span and work



series  $T_{\infty}(A + B) = T_{\infty}(A) + T_{\infty}(B)$


# Composing span and work



series  $T_{\infty}(A + B) = T_{\infty}(A) + T_{\infty}(B)$   
parallel  $T_{\infty}(A \parallel B) = \max(T_{\infty}(A), T_{\infty}(B))$



# Composing span and work



series  $T_{\infty}(A + B) = T_{\infty}(A) + T_{\infty}(B)$

parallel  $T_{\infty}(A \parallel B) = \max(T_{\infty}(A), T_{\infty}(B))$

series or parallel  $T_1 = T_1(A) + T_1(B)$ . **Why?**

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad (\text{I.H.})$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad T(n) \leq a\phi^n - b \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad T(n) \leq a\phi^n - b \quad (\text{I.H.})$$

Substitute the I.H.

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

$$T(n) \leq a(\phi^{n-1} + \phi^{n-2}) - 2b + \Theta(1)$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad T(n) \leq a\phi^n - b \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

Substitute the I.H.

$$\begin{aligned} T(n) &\leq a(\phi^{n-1} + \phi^{n-2}) - 2b + \Theta(1) \\ &= a \frac{\phi + 1}{\phi^2} \phi^n - b + (\Theta(1) - b) \end{aligned}$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad T(n) \leq a\phi^n - b \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

Substitute the I.H.

$$\begin{aligned} T(n) &\leq a(\phi^{n-1} + \phi^{n-2}) - 2b + \Theta(1) \\ &= a \frac{\phi + 1}{\phi^2} \phi^n - b + (\Theta(1) - b) \\ &\leq a \frac{\phi + 1}{\phi^2} \phi^n - b \quad \text{for } b \text{ large} \end{aligned}$$



# Work of Parallel Fibonacci

Write  $T(n)$  for  $T_1$  on input  $n$ .

$$T(n) = T(n-1) + T(n-2) + \Theta(1) \quad T(n) \leq a\phi^n - b \quad (\text{I.H.})$$

Let  $\phi \approx 1.62$  be the solution to

$$\phi^2 = \phi + 1$$

We can show by induction (twice)  
that

$$T(n) \in \Theta(\phi^n)$$

Substitute the I.H.

$$\begin{aligned} T(n) &\leq a(\phi^{n-1} + \phi^{n-2}) - 2b + \Theta(1) \\ &= a \frac{\phi + 1}{\phi^2} \phi^n - b + (\Theta(1) - b) \\ &\leq a \frac{\phi + 1}{\phi^2} \phi^n - b \quad \text{for } b \text{ large} \\ &= a\phi^n - b \end{aligned}$$

# Span and Parallelism of Fibonacci

$$\begin{aligned}T_{\infty}(n) &= \max(T_{\infty}(n-1), T_{\infty}(n-2)) + \Theta(1) \\ &= T_{\infty}(n-1) + \Theta(1)\end{aligned}$$

# Span and Parallelism of Fibonacci

$$\begin{aligned}T_{\infty}(n) &= \max(T_{\infty}(n-1), T_{\infty}(n-2)) + \Theta(1) \\ &= T_{\infty}(n-1) + \Theta(1)\end{aligned}$$

Transforming to sum, we get

$$T_{\infty} \in \Theta(n)$$

# Span and Parallelism of Fibonacci

$$\begin{aligned}T_{\infty}(n) &= \max(T_{\infty}(n-1), T_{\infty}(n-2)) + \Theta(1) \\ &= T_{\infty}(n-1) + \Theta(1)\end{aligned}$$

Transforming to sum, we get

$$T_{\infty} \in \Theta(n)$$

$$\text{parallelism} = \frac{T_1(n)}{T_{\infty}(n)} = \Theta\left(\frac{\phi^n}{n}\right)$$

# Span and Parallelism of Fibonacci

$$\begin{aligned}T_{\infty}(n) &= \max(T_{\infty}(n-1), T_{\infty}(n-2)) + \Theta(1) \\ &= T_{\infty}(n-1) + \Theta(1)\end{aligned}$$

Transforming to sum, we get

$$T_{\infty} \in \Theta(n)$$

$$\text{parallelism} = \frac{T_1(n)}{T_{\infty}(n)} = \Theta\left(\frac{\phi^n}{n}\right)$$

► So an **inefficient** way to compute Fibonacci, but **very parallel**

# Contents

## Dynamic Multithreaded Algorithms

Fork-Join Model

## Using More than One Processor

Capabilities

Span, Work, And Parallelism

Parallel Loops

Scheduling

Race Conditions

# Parallel Loops

```
parallel for  $i = 1$  to  $n$  do  
    statement...  
    statement...  
end for
```

- ▶ Run  $n$  copies in parallel with local setting of  $i$ .

# Parallel Loops

```
parallel for  $i = 1$  to  $n$  do  
    statement...  
    statement...  
end for
```

- ▶ Run  $n$  copies in parallel with local setting of  $i$ .
- ▶ Effectively  $n$ -way spawn



# Parallel Loops

```
parallel for  $i = 1$  to  $n$  do  
    statement...  
    statement...  
end for
```

- ▶ Run  $n$  copies in parallel with local setting of  $i$ .
- ▶ Effectively  $n$ -way spawn
- ▶ Can be implemented with spawn and sync

# Parallel Loops

```
parallel for  $i = 1$  to  $n$  do  
    statement...  
    statement...  
end for
```

- ▶ Run  $n$  copies in parallel with local setting of  $i$ .
- ▶ Effectively  $n$ -way spawn
- ▶ Can be implemented with spawn and sync
- ▶ Span

$$T_{\infty}(n) = \Theta(\log n) + \max_i T_{\infty}(\text{iteration } i)$$

# Parallel Matrix-Vector product

To compute  $y = Ax$ , in parallel

$$y_i = \sum_{j=1}^n a_{ij} x_j$$

# Parallel Matrix-Vector product

To compute  $y = Ax$ , in parallel

$$y_i = \sum_{j=1}^n a_{ij}x_j$$

```
function ROWMULT(A,x,y,i)
   $y_i = 0$ 
  for  $j = 1$  to  $n$  do
     $y_i = y_i + a_{ij}x_j$ 
  end for
end function
```

# Parallel Matrix-Vector product

To compute  $y = Ax$ , in parallel

$$y_i = \sum_{j=1}^n a_{ij}x_j$$

**function** ROWMULT(A,x,y,i)

$$y_i = 0$$

**for**  $j = 1$  to  $n$  **do**

$$y_i = y_i + a_{ij}x_j$$

**end for**

**end function**

**function** MAT-VEC( $A, x, y$ )

Let  $n = \text{rows}(A)$

**parallel for**  $i = 1$  to  $n$  **do**

    RowMult(A,x,y,i)

**end for**

**end function**

# Parallel Matrix-Vector product

To compute  $y = Ax$ , in parallel

$$y_i = \sum_{j=1}^n a_{ij}x_j$$

```
function ROWMULT(A,x,y,i)
     $y_i = 0$ 
    for  $j = 1$  to  $n$  do
         $y_i = y_i + a_{ij}x_j$ 
    end for
end function
```

```
function MAT-VEC( $A, x, y$ )
    Let  $n = \text{rows}(A)$ 
    parallel for  $i = 1$  to  $n$  do
        RowMult( $A, x, y, i$ )
    end for
end function
```

$$T_1(n) \in \Theta(n^2) \quad (\text{serialization})$$
$$T_\infty(n) = \underbrace{\Theta(\log(n))}_{\text{parallel for}} + \underbrace{\Theta(n)}_{\text{RowMult}}$$

# Parallel Matrix-Vector product

```
function ROWMULT(A,x,y,i)
     $y_i = 0$ 
    for  $j = 1$  to  $n$  do
         $y_i = y_i + a_{ij}x_j$ 
    end for
end function
```

```
function MAT-VEC( $A, x, y$ )
    Let  $n = \text{rows}(A)$ 
    parallel for  $i = 1$  to  $n$  do
        RowMult( $A, x, y, i$ )
    end for
end function
```

# Parallel Matrix-Vector product

```
function ROWMULT(A,x,y,i)
     $y_i = 0$ 
    for  $j = 1$  to  $n$  do
         $y_i = y_i + a_{ij}x_j$ 
    end for
end function
```


```
function MAT-VEC( $A, x, y$ )
    Let  $n = \text{rows}(A)$ 
    parallel for  $i = 1$  to  $n$  do
        RowMult( $A, x, y, i$ )
    end for
end function
```

► Why is RowMult not using parallel for?



# Divide and Conquer Matrix-Vector product

```
function MVDC( $A, x, y, f, t$ )  
  if  $f == t$  then  
    RowMult( $A, x, y, f$ )  
  else  
     $m = \lfloor (f + t) / 2 \rfloor$   
    spawn MVDC( $A, x, y, f, m$ )  
    MVDC( $A, x, y, m + 1, t$ )  
    sync  
  end if  
end function
```



# Divide and Conquer Matrix-Vector product

```
function MVDC( $A, x, y, f, t$ )  
  if  $f == t$  then  
    RowMult( $A, x, y, f$ )  
  else  
     $m = \lfloor (f + t) / 2 \rfloor$   
    spawn MVDC( $A, x, y, f, m$ )  
    MVDC( $A, x, y, m + 1, t$ )  
  sync  
end if  
end function
```

►  $T_{\infty}(n) = \Theta(\log n)$   
(binary tree) +  
 $T_{\infty}(\text{RowMult})$

# Divide and Conquer Matrix-Vector product

```
function MVDC( $A, x, y, f, t$ )  
  if  $f == t$  then  
    RowMult( $A, x, y, f$ )  
  else  
     $m = \lfloor (f + t) / 2 \rfloor$   
    spawn MVDC( $A, x, y, f, m$ )  
    MVDC( $A, x, y, m + 1, t$ )  
  sync  
end if  
end function
```

- ▶  $T_{\infty}(n) = \Theta(\log n)$   
(binary tree) +  
 $T_{\infty}(\text{RowMult})$
- ▶  $\Theta(n)$  leaves (one per row)

# Divide and Conquer Matrix-Vector product

```
function MVDC( $A, x, y, f, t$ )  
  if  $f == t$  then  
    RowMult( $A, x, y, f$ )  
  else  
     $m = \lfloor (f + t) / 2 \rfloor$   
    spawn MVDC( $A, x, y, f, m$ )  
    MVDC( $A, x, y, m + 1, t$ )  
    sync  
  end if  
end function
```

- ▶  $T_{\infty}(n) = \Theta(\log n)$   
(binary tree) +  
 $T_{\infty}(\text{RowMult})$
- ▶  $\Theta(n)$  leaves (one per row)
- ▶  $\Theta(n)$  interior nodes  
(binary tree)

# Divide and Conquer Matrix-Vector product

```
function MVDC( $A, x, y, f, t$ )  
  if  $f == t$  then  
    RowMult( $A, x, y, f$ )  
  else  
     $m = \lfloor (f + t) / 2 \rfloor$   
    spawn MVDC( $A, x, y, f, m$ )  
    MVDC( $A, x, y, m + 1, t$ )  
    sync  
  end if  
end function
```

- ▶  $T_{\infty}(n) = \Theta(\log n)$   
(binary tree) +  $T_{\infty}(\text{RowMult})$
- ▶  $\Theta(n)$  leaves (one per row)
- ▶  $\Theta(n)$  interior nodes (binary tree)
- ▶  $T_1(n) = \Theta(n^2)$

# Contents

## Dynamic Multithreaded Algorithms

Fork-Join Model

## Using More than One Processor

Capabilities

Span, Work, And Parallelism

Parallel Loops

Scheduling

Race Conditions

# Scheduling

## Scheduling Problem

**Abstractly** Mapping threads to processors

**Pragmatically** Mapping logical threads to a thread pool.

# Scheduling

## Scheduling Problem

**Abstractly** Mapping threads to processors

**Pragmatically** Mapping logical threads to a thread pool.

## Ideal Scheduler

**On-Line** No advance knowledge of when threads will spawn or complete.

**Distributed** No central controller.



# Scheduling

## Scheduling Problem

**Abstractly** Mapping threads to processors

**Pragmatically** Mapping logical threads to a thread pool.

## Ideal Scheduler

**On-Line** No advance knowledge of when threads will spawn or complete.

**Distributed** No central controller.

► to simplify analysis, we relax the second condition

# A greedy centralized scheduler

Maintain a *ready queue* of strands ready to run.

## Scheduling Step

**Complete Step** If  $\geq p$  ( $\#$  processors) strands are ready, assign  $p$  strands to processors.

# A greedy centralized scheduler

Maintain a *ready queue* of strands ready to run.

## Scheduling Step

**Complete Step** If  $\geq p$  ( $\#$  processors) strands are ready, assign  $p$  strands to processors.

**Incomplete Step** Otherwise, assign all waiting strands to processors

# A greedy centralized scheduler

Maintain a *ready queue* of strands ready to run.

## Scheduling Step

**Complete Step** If  $\geq p$  ( $\#$  processors) strands are ready, assign  $p$  strands to processors.

**Incomplete Step** Otherwise, assign all waiting strands to processors

- To simplify analysis, split any non-unit strands into a chain of unit strands

# A greedy centralized scheduler

Maintain a *ready queue* of strands ready to run.

## Scheduling Step

**Complete Step** If  $\geq p$  ( $\#$  processors) strands are ready, assign  $p$  strands to processors.

**Incomplete Step** Otherwise, assign all waiting strands to processors

- ▶ To simplify analysis, split any non-unit strands into a chain of unit strands
- ▶ Therefore, after one time step, we schedule again.

# Optimal and Approximate Scheduling

Recall

$$T_p \geq T_1/p \quad (\text{work law})$$

$$T_p \geq T_\infty \quad (\text{span})$$

Therefore

$$T_p \geq \max(T_1/p, T_\infty) = \text{opt}$$

# Optimal and Approximate Scheduling

Recall

$$T_p \geq T_1/p \quad (\text{work law})$$

$$T_p \geq T_\infty \quad (\text{span})$$

Therefore

$$T_p \geq \max(T_1/p, T_\infty) = \text{opt}$$

With the greedy algorithm we can achieve

$$T_p \leq \frac{T_1}{p} + T_\infty \leq 2 \max(T_1/p, T_\infty) = 2 \times \text{opt}$$

# Counting Complete Steps

- ▶ Let  $k$  be the number of complete steps.



# Counting Complete Steps

- ▶ Let  $k$  be the number of complete steps.
- ▶ At each complete step we do  $p$  units of work.


# Counting Complete Steps

- ▶ Let  $k$  be the number of complete steps.
- ▶ At each complete step we do  $p$  units of work.
- ▶ Every unit of work corresponds to one step of the serialization, so  $kp \leq T_1$ .

# Counting Complete Steps


- ▶ Let  $k$  be the number of complete steps.
- ▶ At each complete step we do  $p$  units of work.
- ▶ Every unit of work corresponds to one step of the serialization, so  $kp \leq T_1$ .
- ▶ Therefore  $k \leq T_1/p$

# Counting Incomplete Steps




► Let  $G$  be the DAG of *remaining strands*.

# Counting Incomplete Steps




- ▶ Let  $G$  be the DAG of *remaining strands*.
- ▶ The **ready queue** of strands is exactly the set of sources in  $G$

# Counting Incomplete Steps




- ▶ Let  $G$  be the DAG of *remaining strands*.
- ▶ The **ready queue** of strands is exactly the set of sources in  $G$
- ▶ In incomplete step runs *all* sources in  $G$

# Counting Incomplete Steps



- ▶ Let  $G$  be the DAG of *remaining strands*.
- ▶ The **ready queue** of strands is exactly the set of sources in  $G$
- ▶ In incomplete step runs *all* sources in  $G$
- ▶ Every longest path starts at a source (otherwise, extend)


# Counting Incomplete Steps



- ▶ Let  $G$  be the DAG of *remaining strands*.
- ▶ The **ready queue** of strands is exactly the set of sources in  $G$
- ▶ In incomplete step runs *all* sources in  $G$
- ▶ Every longest path starts at a source (otherwise, extend)
- ▶ After an incomplete step, length of longest path shrinks by 1



# Counting Incomplete Steps



- ▶ Let  $G$  be the DAG of *remaining strands*.
- ▶ The **ready queue** of strands is exactly the set of sources in  $G$
- ▶ In incomplete step runs *all* sources in  $G$
- ▶ Every longest path starts at a source (otherwise, extend)
- ▶ After an incomplete step, length of longest path shrinks by 1
- ▶ There can be at most  $T_\infty$  steps.

# Parallel Slackness

$$\text{parallel slackness} = \frac{\text{parallelism}}{p} = \frac{T_1}{pT_\infty}$$

- ▶ If slackness  $< 1$ , speedup  $< p$

# Parallel Slackness

$$\text{parallel slackness} = \frac{\text{parallelism}}{p} = \frac{T_1}{pT_\infty}$$

$$\text{speedup} = \frac{T_1}{T_p} \leq \frac{T_1}{T_\infty} = p \times \text{slackness}$$

- ▶ If slackness  $< 1$ , speedup  $< p$
- ▶ If slackness  $\geq 1$ , linear speedup achievable for given number of processors

# Slackness and Scheduling

$$\text{slackness} := \frac{T_1}{p \times T_\infty}$$

## Theorem

*For sufficiently large slackness, the greed scheduler approaches time  $T_1/p$ .*

# Slackness and Scheduling

$$\text{slackness} := \frac{T_1}{p \times T_\infty}$$

## Theorem

*For sufficiently large slackness, the greed scheduler approaches time  $T_1/p$ .*

Suppose

$$\frac{T_1}{p \times T_\infty} \geq c$$

# Slackness and Scheduling

$$\text{slackness} := \frac{T_1}{p \times T_\infty}$$

Then

$$T_\infty \leq \frac{T_1}{cp} \quad (2)$$

## Theorem

*For sufficiently large slackness, the greed scheduler approaches time  $T_1/p$ .*

Suppose

$$\frac{T_1}{p \times T_\infty} \geq c$$

# Slackness and Scheduling

$$\text{slackness} := \frac{T_1}{p \times T_\infty}$$

## Theorem

*For sufficiently large slackness, the greedy scheduler approaches time  $T_1/p$ .*

Suppose

$$\frac{T_1}{p \times T_\infty} \geq c$$

Then

$$T_\infty \leq \frac{T_1}{cp} \quad (2)$$

Recall that with the greedy scheduler,

$$T_p \leq \left( \frac{T_1}{p} + T_\infty \right)$$

# Slackness and Scheduling

$$\text{slackness} := \frac{T_1}{p \times T_\infty}$$

## Theorem

*For sufficiently large slackness, the greedy scheduler approaches time  $T_1/p$ .*

Suppose

$$\frac{T_1}{p \times T_\infty} \geq c$$

Then

$$T_\infty \leq \frac{T_1}{cp} \quad (2)$$

Recall that with the greedy scheduler,

$$T_p \leq \left( \frac{T_1}{p} + T_\infty \right)$$

Substituting (2), we have

$$T_p \leq \frac{T_1}{p} \left( 1 + \frac{1}{c} \right)$$



# Contents

## Dynamic Multithreaded Algorithms

Fork-Join Model

## Using More than One Processor

Capabilities

Span, Work, And Parallelism

Parallel Loops

Scheduling

Race Conditions

# Race Conditions

## Non-Determinism

- ▶ result varies from run to run
- ▶ sometimes OK (in certain randomized algorithms)
- ▶ mostly a bug.

# Race Conditions


## Non-Determinism

- ▶ result varies from run to run
- ▶ sometimes OK (in certain randomized algorithms)
- ▶ mostly a bug.

## Example


```
x = 0
parallel for i ← 1 to 2 do
  x ← x + 1
```

# Racy execution




- ▶ all possible topological sorts are valid execution orders

# Racy execution




- ▶ all possible topological sorts are valid execution orders
- ▶ In particular it's not hard for both loads to complete before either store

# Racy execution



- ▶ all possible topological sorts are valid execution orders
- ▶ In particular it's not hard for both loads to complete before either store
- ▶ In practice there are various synchronization strategies (locks, etc...).

# Racy execution



- ▶ all possible topological sorts are valid execution orders
- ▶ In particular it's not hard for both loads to complete before either store
- ▶ In practice there are various synchronization strategies (locks, etc...).
- ▶ Here we will insist that parallel strands are **independent**

# We can write bad code with spawn too

```
sum(i, j)
  if (i>j)
    return;
  if (i==j)
    x++;
  else
    m=(i+j)/2;
    spawn sum(i,m);
    sum(m+1,j);
  sync;
```

- ▶ here we have the same non-deterministic interleaving of reading and writing  $x$
- ▶ the style is a bit unnatural, in particular we are not using the return value of spawn at all.



# Being more *functional* helps

```
sum(i, j)
  if (i>j) return 0;
  if (i==j) return 1;
```

```
m ← (i+j)/2;
```

```
left ← spawn sum(i,m);
right ← sum(m+1,j);
sync;
return left + right;
```

► each strand writes into different variables

# Being more *functional* helps

```
sum(i, j)
  if (i>j) return 0;
  if (i==j) return 1;

  m ← (i+j)/2;

  left ← spawn sum(i,m);
  right ← sum(m+1,j);
  sync;
  return left + right;
```

- ▶ each strand writes into different variables
- ▶ sync is used as a **barrier** to serialize

# Single Writer races

```
x ← spawn foo(x)
y ← foo(x)
sync
```

- ▶ arguments to spawned routines are evaluated in the parent context

# Single Writer races

```
x ← spawn foo(x)
y ← foo(x)
sync
```

- ▶ arguments to spawned routines are evaluated in the parent context
- ▶ but this isn't enough to be race free.

# Single Writer races

```
x ← spawn foo(x)
y ← foo(x)
sync
```

- ▶ arguments to spawned routines are evaluated in the parent context
- ▶ but this isn't enough to be race free.
- ▶ which value  $x$  is passed to the second call of 'foo' depends how long the first one takes.